A Consensus-Trained U-Net with Learned Instance Selection on a Threshold Tree

A Solution to the Solar Filament Segmentation Challenge 2026

Mukharbek Organokov
mukharbek.organokov@gmail.com

ABSTRACT

This report describes our solution to the Solar Filament Segmentation Challenge 2026, a Kaggle competition on automatic segmentation of solar filaments in GONG H-α observations. A single U-Net with a ResNet-34 encoder, run at 1536 px, is trained on the consensus of the annotators and averaged over four flipped views. Its connected components at twelve thresholds form a tree of nested candidates; a gradient-boosted hurdle model, trained on out-of-fold candidates, predicts each one's expected contribution to Panoptic Quality (PQ) from its shape, contrast, position, place in the tree and the disagreement of the four views, and the disjoint set with the largest expected gain is chosen exactly on the tree. A second stage redraws each chosen filament at native resolution in a crop turned along its own axis, with ten refiners (five UNet++ and five U-Nets) averaged. On all 707 training images, scored out of fold with a split by calendar month, the pipeline reaches PQ 0.440 (95% bootstrap interval 0.431–0.449), against 0.427 without the second stage, 0.422 with candidates from a single threshold and 0.409 for the same network at 1024 px with a hand-tuned rule; one MAGFiLO annotator scored against another on the same images reaches 0.340. The public leaderboard gives 0.37 for this pipeline and for its first version of the second stage, 0.36 for the first stage alone and 0.35 for the earlier pipelines. We list the variants that did not help and give evidence that most of the remaining loss is disagreement between annotators.

1 INTRODUCTION

This report goes over the details of our submitted solution to the Solar Filament Segmentation Challenge 2026 [2]. The challenge, launched on July 10, 2026, is organized by the Earth-Space AI Research (ESAIR) Lab[1]. This challenge utilizes the MAGFiLO dataset for evaluation of the segmentation algorithms [1].

One property of the data shapes every decision: the annotators disagree strongly, and each annotator is scored separately against the same prediction. A prediction can therefore never match all annotators, and the best achievable PQ is far below 1. Our solution is (i) a simple, reproducible segmenter trained on the annotators' consensus; (ii) learned instance selection: candidates from twelve thresholds, a scorer of their expected PQ contribution and an exact choice of the best disjoint set; and (iii) an error analysis, with the annotator-vs-annotator score as a reference, that explains why the usual improvements did not help. Only the competition's training images and annotations are used; no external filament masks, spines or chirality labels. Code, weights and a notebook that reproduces the submission are in the repository, and every experiment is logged in `docs/DESIGN.md`.

[1] www.esairlab.com/

Kaggle Competition, 2026, kaggle.com/competitions/filament-segmentation-2026
2026.

2 METHODOLOGY

2.1 Data and metric

The training set has 707 images (2048×2048 px) with 8,199 annotated filaments; 411 images have one annotator, 145 two and 151 three, giving 1,154 annotator-image entries. The test set has 180 images. Within each entry, a ground-truth (GT) and a predicted mask form a true positive (TP) when their IoU exceeds 0.5, and PQ [7] is $\sum_{TP} \mathrm{IoU}/(|TP| + 0.5|FP| + 0.5|FN|)$, pooled over all entries. It factors into segmentation quality SQ (the mean IoU of the TPs) and recognition quality RQ. Our implementation reproduces the organisers' self-evaluation notebook exactly. That code does not enforce one-to-one matching, so overlapping predictions of one filament would each count as a TP; all our submissions are disjoint by construction.

2.2 Input, target and network

Each image is area-downsampled from 2048 to 1536 px and normalised with fixed constants; there is no other pre-processing. The training target is the consensus, i.e. the fraction of annotators marking each pixel, computed at 2048 px and then downsampled. With binary cross-entropy (BCE) on this soft target the network learns the probability that an annotator marks a pixel, which is the quantity the multi-annotator metric rewards. The network is a U-Net [10] from `segmentation_models_pytorch` [6] with an ImageNet ResNet-34 encoder [5] adapted to one input channel (24.4 M parameters). Loss: BCE + soft Dice. AdamW, learning rate $5 \cdot 10^{-4}$ with one epoch of warm-up and cosine decay, weight decay 10^{-4}, batch 4, 40 epochs. Augmentation: flips, 90-degree rotations, affine (±30°, scale 0.9–1.1), brightness, contrast and gamma. The submitted model is trained on all 707 images (36 min on one NVIDIA A100, about 3 h on an Apple M3 Max).

2.3 Candidates: a tree over thresholds

At inference the network sees four views (identity, horizontal, vertical and both flips); their mean is the probability map and their per-pixel standard deviation (the *view spread*) marks where the model is unsure. Both are upsampled bilinearly to 2048 px. The map is thresholded at twelve levels, 0.25 to 0.80 in steps of 0.05. At each level a morphological closing with a 31 px disk decides which fragments belong to one filament (the masks keep the thresholded pixels), and the 8-connected components of at least 100 px are candidates. Thresholding and closing are monotone, so every component at a higher level lies inside exactly one component at the next

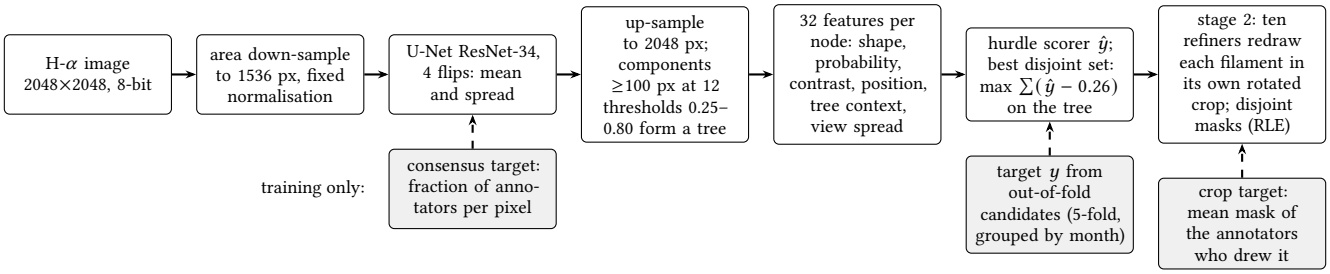


Figure 1: The end-to-end pipeline. White: inference, applied to every test image. Grey: what each learned stage is trained on.

lower level: the candidates form a forest. A filament can thus be taken whole at a low threshold, thinner at a high one, or separated at a higher level from a neighbour it merges with lower down (Fig. 1).

2.4 Instance selection

Every false positive costs half a point in the PQ denominator, so which candidates are submitted matters as much as their shape. A kept candidate that matches k of the n entries of its image adds about $k \cdot \text{IoU}$ to the numerator and $0.5n$ to the denominator either way, so it pays when its expected matched IoU per entry, $y = \frac{1}{n}\sum_e \text{IoU}_e\,[\text{IoU}_e > 0.5]$ over the entries e of its image, exceeds about PQ/2. We predict y for every node of the tree with a hurdle model, $\hat{y} = P(y > 0) \cdot E[y \mid y > 0]$, both parts gradient-boosted trees from scikit-learn [8] (`HistGradientBoosting`, 400 iterations, learning rate 0.05, 31 leaves, at least 40 samples per leaf, L2 1.0), trained on out-of-fold nodes with rows weighted by the image's number of entries, as pooled PQ weights them. The 32 features are: area and log-area; the mean, maximum, standard deviation, 10/50/90% quantiles of the probability inside and the fraction above 0.8; elongation, length, width and compactness; the brightness inside and its contrast with a 15 px ring; distance from the disk centre, latitude and longitude on a disk fitted with a Hough circle; the number of candidates and the mean probability of the image; the node's threshold, its area relative to its parent, its number of children and siblings, the share of its area its children keep, and the number of levels over which it keeps 85% of its area; and the mean, 90% quantile, maximum and relative value of the view spread inside it. The submission is the set of disjoint nodes (an antichain of the forest) maximising $\sum(\hat{y} - \tau)$, which dynamic programming solves exactly: $f(v) = \max(\hat{y}_v - \tau, \sum_c f(c))$ over the children c of v. We use $\tau = 0.26$, the middle of a range (0.24–0.28) over which OOF PQ varies by 0.0004.

2.5 Stage 2: refinement in rotated crops

A level set of one map cannot draw every filament: a selected node may be a fragment, or thinner or wider than the annotators drew it. Following the second stage of a public notebook [11], each selected filament (the *seed*) is redrawn in a 384×128 px crop centred on it, its principal axis turned horizontal, covering 1.5 times its length at native resolution (scaled down only when it does not fit). The four input channels are the image, the stage-1 probability, the seed and the other seeds of the image. The target is the mean of the masks drawn for this filament by the annotators who drew it (each entry's mask that overlaps the seed most, if its IoU exceeds 0.3); whether the filament exists stays stage 1's decision. Two refiners with a ResNet-34 encoder are trained per fold on the other folds' crops (BCE + soft Dice): a UNet++ [12] for 30 epochs on the 4,429 seed crops, and a U-Net for 15 epochs whose training set also holds 6,610 crops of tree nodes at other thresholds (fragments and seeds cut at other levels, kept when they match an annotation at IoU > 0.3 or lie at least half inside it), with per-sample flips and a small affine jitter. At test time the ten models (five folds, two architectures) and four flips are averaged. The output is thresholded at 0.40, fixed before the evaluation, and pasted back: only pieces touching the seed are kept, a seed covered at least half by its neighbours' new masks is absorbed, and overlaps go to the higher refined probability, so the masks stay disjoint. A seed whose crop would need scaling below 0.5, or that the refiner erases, keeps its stage-1 mask. Inference with both stages takes about 16 s per image on a shared M3 Max.

2.6 Validation split and tuning

Filaments persist for days, so images from the same week are near-duplicates. We use a 5-fold `GroupKFold` grouped by calendar month (YYYYMM), so no month is in both training and validation. The threshold and closing are grid-searched with the exact metric on the pooled out-of-fold (OOF) probabilities of the five fold models. The scorer is cross-validated on the same folds: the scorer applied to a fold's nodes is trained on the other four, so every OOF number below comes from models that never saw the image. The view spread of the OOF images is computed by each fold's model, as at test time. All choices were made on OOF, never on the public leaderboard, which covers about 90 test images and shows two decimals. The scorer threshold τ, the stage-1 settings and the choice among five refiner blends were made on the same OOF folds, so the OOF figures carry some selection optimism; the stage-2 pixel threshold was fixed before its evaluation.

3 EVALUATION

3.1 Main results

Table 1 gives the main numbers. Uncertainty comes from two sources: the spread of PQ over the five folds, and a bootstrap that resamples whole images (all entries of an image together, 2,000 resamples) and recomputes the pooled PQ. The paired bootstrap of the difference to the previous pipeline resamples the same images for both. To put the numbers in scale, each annotator's masks are

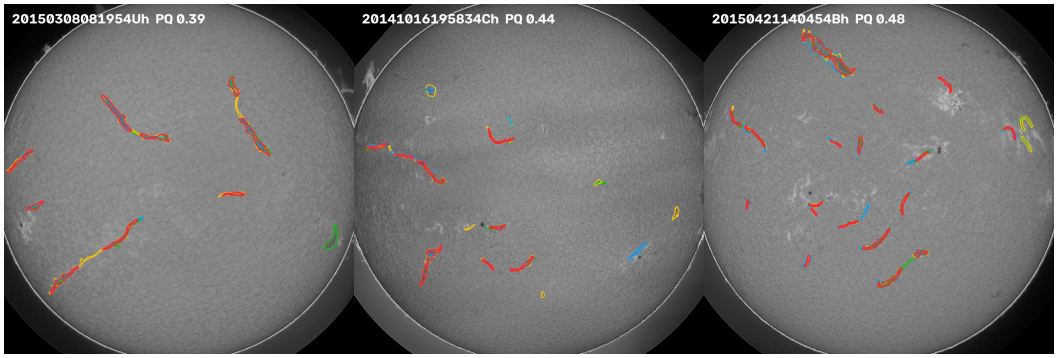


Figure 2: Out-of-fold predictions of stage 1 (red) on three validation images with three annotators (green, yellow, blue), at the 50th, 75th and 90th percentile of image PQ among images with at least five filaments per annotator. Many misses are filaments marked by one annotator only.

Table 1: Main results. OOF: out of fold, 707 images, 1,154 entries. Brackets: 95% bootstrap interval over images.

	PQ	SQ	RQ
Final (stage 2), OOF	**0.440** [0.431, 0.449]	0.680	0.647
per fold, mean ± s.d.	0.440 ± 0.009		
296 multi-annotator images	0.441	0.677	0.651
Annotator vs. annotator (same 296)	0.340	0.634	0.536
1024 px, 2 views, hand rule	0.409 [0.399, 0.418]	0.666	0.614
+ 4 views	0.412	0.667	0.617
+ component scorer	0.417	0.668	0.624
1536 px, 4 views, hand rule	0.417	0.671	0.621
+ scorer, one threshold	0.422 [0.412, 0.431]	0.674	0.626
+ tree (6 levels), hurdle	0.426	0.676	0.630
+ 12 levels, spread (stage 1)	0.427 [0.417, 0.436]	0.675	0.632
+ stage 2, one U-Net refiner	0.433	0.677	0.639
+ UNet++ and extra-crop refiners (final)	0.440	0.680	0.647
Stage 2 over stage 1 (paired)	+0.013 [0.010, 0.016]		
Final over one refiner (paired)	+0.007 [0.005, 0.009]		
Stage 1 over one threshold (paired)	+0.005 [0.001, 0.008]		
Stage 1 over 1024 px rule (paired)	+0.018 [0.013, 0.023]		
Public LB, stage 2, one refiner	0.37	–	–
Public LB, final (ten refiners)	0.37	–	–
Public LB, stage 1	0.36	–	–
Public LB, earlier pipelines	0.35	–	–

scored as predictions against every other annotator of the same image (296 images with two or three annotators).

The gains accumulate in steps that each hold up in a paired bootstrap. Resolution and views: tuned per fold, 1536 px beats 1024 px on four of five folds (+0.009, −0.002, +0.009, +0.005, +0.007), and 2048 px matches 1536 px on the two folds tried. Scoring: with a single threshold the scorer beats the submitted hand rule (confidence ≥ 0.8) on every fold, by 0.002–0.009, but the best pooled hand rule (confidence ≥ 0.85, PQ 0.420) by only 0.002. Tree: candidates from six thresholds with the hurdle model and the context features give +0.004 [0.001, 0.008]; twelve thresholds alone add nothing (0.425), and the view-spread features bring the gain to +0.005 [0.001, 0.008] over the single threshold, on four of five folds (per fold 0.434, 0.425, 0.411, 0.430, 0.434 against 0.432, 0.415, 0.407, 0.421, 0.435). Against the 1024 px rule stage 1 gains +0.018 [0.013, 0.023], on every fold. Stage 2: each fold's crops are redrawn by the refiners that did not see them. One U-Net refiner raises PQ on all five folds (0.434, 0.433, 0.419, 0.442, 0.437), by +0.006 [0.003, 0.009] pooled, through RQ (0.632 to 0.639) more than SQ (0.675 to 0.677): redrawn fragments now cross IoU 0.5. Against it, a UNet++ gains +0.002 [0.000, 0.005], the U-Net with the extra crops and augmentation +0.005 [0.003, 0.008], and the average of these two +0.007 [0.005, 0.009], each on all five folds; the final per-fold PQ is 0.442, 0.438, 0.427, 0.448, 0.446. The blend was the best of five on every fold, but its margin over the extra-crop U-Net alone is +0.0016 [−0.0002, 0.0034], inside noise. The pixel threshold barely matters (within 0.0011 from 0.35 to 0.45 for each refiner). The leaderboard moved by one step at its two decimals for the first stage and for the first refiner (0.35 to 0.36 to 0.37) and did not move for the ten-refiner version, whose expected gain (+0.007) is a quarter of a step. This is consistent with, not evidence for, the gains: scaled to the 90 or so public images, the bootstrap interval of a single pipeline is about ±0.027, larger than any of them. Local OOF ran 0.06–0.07 above the public leaderboard for every submission (0.07 for the final one). The two public OOF-validated notebooks we read report a gap of 0.05–0.06 and attribute it to the labels (sparser test annotations); one of them finds the test frames indistinguishable from training frames. We have not tested this; a label-side gap fits the annotator disagreement below, and our ranking rests on paired OOF differences, not on the leaderboard.

Fig. 3 shows the three evaluation plots requested by the organisers on the OOF predictions of stage 1. Over GT/prediction pairs with IoU > 0, the mean IoU is 0.575 and the mean Dice 0.708 (0.553 and 0.688 for the 1024 px rule pipeline). 3.1% of predicted filaments overlap more than one annotation (merges) and 3.8% of annotations overlap more than one prediction (splits); 10.6% of predictions overlap no annotation of their entry and 22.7% of annotations overlap no prediction. Fig. 2 shows typical images.

3.2 What did not help

Table 2 lists the variants we tried, each with its own tuned post-processing. A Mask R-CNN [4] detector learns grouping explicitly, but its masks are coarser (SQ 0.642 vs. 0.657 for a U-Net trained

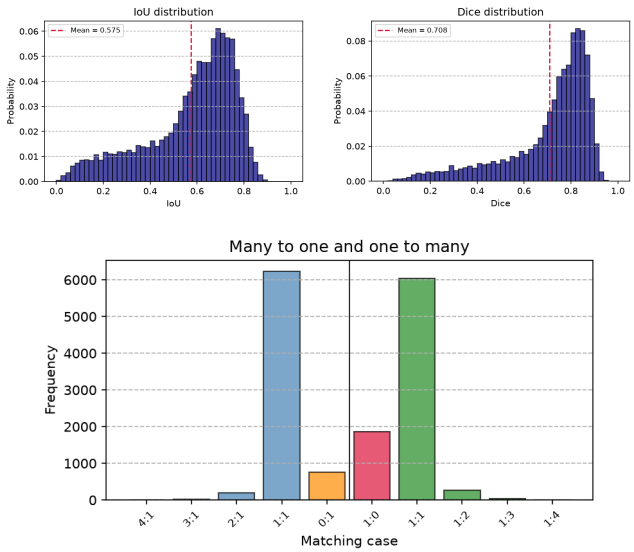


Figure 3: Out-of-fold IoU and Dice distributions over GT/prediction pairs with IoU > 0, and the one-to-many / many-to-one counts (left: number of GT linked to each prediction; right: number of predictions linked to each GT).

for a similar number of epochs); adding its extra instances, or an axis-aligned crop refiner around each instance following Reche & Cid [9], changes PQ by less than 0.002 on the fully trained U-Net. A ResNet-50 encoder, an input divided by its own radial brightness profile (limb-darkening correction) and a Lovász hinge loss [3] match or underperform the baseline; with the corrected input the best threshold moves from 0.45 to 0.65, so the network becomes more confident but not more accurate. Averaging two equal-strength models with different encoders gives the PQ of the better one (a ConvNeXt-T U-Net averaged with the ResNet-34 gains 0.0016 on fold 0, below what one fold resolves), and the 5-fold average gives the same leaderboard score as one model: the models share their errors because the errors come from the annotators, so stage 1 is one network. After stage 2, a keep/drop rescoring from the refiners' output gains at most 0.0014 PQ out of fold and isotonic calibration of the scorer nothing, so selection is not adopted as a further lever; the earlier U-Net refiner adds nothing once the extra-crop U-Net is in the blend.

3.3 Where the remaining PQ is lost

After stage 2. Measured on every third OOF image (236 images, with the one-refiner pipeline, PQ 0.426 on this subset against 0.433 on all images; intervals about ±0.007): 1,629 TP, 769 FP and 1,141 FN. Of the FN, 661 overlap no prediction, 139 overlap one at IoU $<$ 0.3 and 341 at IoU 0.3–0.5; the 341 near misses are the same before and after stage 2, while stage 2 adds TP (1,619 to 1,629) and IoU. 24% of predicted filaments match no annotator. Recall by area is 0.16 below 300 px, 0.39 at 300–500 px, 0.57 at 500–1,000 px and 0.70 at 1,000–2,000 px, and annotators rarely agree on small filaments: the share of other annotators marking the same filament (IoU $>$ 0.5) is 0.28 below 300 px and about 0.60 above 1,000 px. An oracle that drops every net-negative predicted filament would reach 0.474 on the subset; the refiners' output does not predict which these are. *Extent.* In the 1024 px rule pipeline, 37% of the 3,174 OOF false negatives were near misses; the prediction was smaller than the GT in 46% of them and larger in 34%, and only 17% were merges and 16% splits. Dilating or eroding every mask by one pixel costs 0.01–0.03 PQ: extent errors are per instance, not a bias a rule could correct. *When a doubtful prediction pays.* A prediction matching k of n annotators raises PQ only if $k/n > \mathrm{PQ}/(2\,\mathrm{SQ}) \approx 0.32$ at the final PQ and SQ. A filament marked by one annotator of three is break-even, which is why lowering the area or confidence thresholds never helped; in the 1024 px rule pipeline 28% of the false positives matched another annotator's filament on the same image.

Table 2: Variants (fold 0, 142 images, tuned PQ) and post-processing rules (5-fold OOF, 707 images).

Variant (fold 0)	PQ
U-Net ResNet-34, 40 epochs (baseline)	**0.416**
10 / 20 epochs	0.389 / 0.396
+ axis-aligned crop refiner	0.413
+ Mask R-CNN instances not covered	0.417
Mask R-CNN alone (disjoint masks)	0.377
ResNet-50 encoder	0.414
Limb-darkening-corrected input	0.411
BCE + Lovász hinge	0.390
Average of ResNet-34 and ResNet-50	0.416
8 dihedral views (4 flips: 0.419)	0.419
2048 px, folds 0 / 3 (1536 px: 0.403 / 0.387)	0.404 / 0.388
Post-processing rule (OOF, 1024 px)	**PQ**
Submission setting (confidence ≥ 0.8)	0.409
Confidence ≥ 0.85	0.411
No closing, no confidence or limb filter	0.398
Size-dependent confidence threshold	0.411
Hysteresis (seed and grow thresholds)	0.409
Masks dilated / eroded by 1 px	0.381 / 0.399
70-degree limb cut-off	no change

3.4 Strengths, weaknesses and possible improvements

Strengths. One network, two small tree models and ten crop refiners, no external data, about 16 s per image, and a submission reproduced byte for byte by the notebook. Selection optimises the metric directly, and the exact antichain keeps the masks disjoint without hand-tuned rules. *Weaknesses.* Missed annotations are the largest error (23% of annotations overlap no prediction), mostly small filaments and filaments marked by one annotator; extent errors of one or two pixels cost SQ; and every gain since the 1024 px rule is below the resolution of the public leaderboard, so the choice rests on paired out-of-fold comparisons. *Improvements.* More diverse members for the map and the refiner.

4 ACKNOWLEDGMENT

The organization of the Filament Segmentation Challenge 2026 is partially supported by the U.S. National Science Foundation (NSF) under Grant No. 2209912 and 2433781, directorate for Computer

and Information Science and Engineering (CSE), and Office of Advanced Cyberinfrastructure (OAC), and Grant No. 2511630, AST Division Of Astronomical Sciences and MPS Directorate for Mathematical and Physical Sciences.

The Filament Segmentation Challenge 2026 is sponsored by the U.S. National Science Foundation (NSF) National Solar Observatory (NSO).

This work utilizes GONG data obtained by the NSO Integrated Synoptic Program, managed by the National Solar Observatory, which is operated by the Association of Universities for Research in Astronomy (AURA), Inc. under a cooperative agreement with the National Science Foundation and with contribution from the National Oceanic and Atmospheric Administration. The GONG network of instruments is hosted by the Big Bear Solar Observatory, High Altitude Observatory, Learmonth Solar Observatory, Udaipur Solar Observatory, Instituto de Astrofísica de Canarias, and Cerro Tololo Interamerican Observatory.

REFERENCES

[1] Azim Ahmadzadeh, Rohan Adhyapak, Kartik Chaurasiya, Laxmi Alekhya Nagubandi, V. Aparna, Petrus C. Martens, Alexei Pevtsov, Luca Bertello, Alexander Pevtsov, Naomi Douglas, Samuel McDonald, Apaar Bawa, Eugene Kang, Riley Wu, Dustin J. Kempton, Aya Abdelkarem, Patrick M. Copeland, and Sri Harsha Seelamneni. 2024. A dataset of manually annotated filaments from H-alpha observations. In *Nature Scientific Data (Nature Scientific Data, Vol. 11)*. Springer, Article 1031, 1031 pages. https://doi.org/10.1038/s41597-024-03876-y

[2] Azim Ahmadzadeh, Dustin J. Kempton, Qin Li, and Alexei A. Pevtsov. 2026. Solar Filament Segmentation Challenge 2026. https://kaggle.com/competitions/filament-segmentation-2026. Kaggle.

[3] Maxim Berman, Amal Rannen Triki, and Matthew B. Blaschko. 2018. The Lovász-Softmax Loss: A Tractable Surrogate for the Optimization of the Intersection-Over-Union Measure in Neural Networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 4413–4421.

[4] Kaiming He, Georgia Gkioxari, Piotr Dollár, and Ross Girshick. 2017. Mask R-CNN. In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*. 2961–2969.

[5] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016. Deep Residual Learning for Image Recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 770–778.

[6] Pavel Iakubovskii. 2019. Segmentation Models Pytorch. https://github.com/qubvel-org/segmentation_models.pytorch.

[7] Alexander Kirillov, Kaiming He, Ross Girshick, Carsten Rother, and Piotr Dollár. 2019. Panoptic Segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 9404–9413.

[8] Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, Jake Vanderplas, Alexandre Passos, David Cournapeau, Matthieu Brucher, Matthieu Perrot, and Édouard Duchesnay. 2011. Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research* 12 (2011), 2825–2830.

[9] Antonio Reche and Consuelo Cid. 2026. New dataset and framework for filament detection, classification, and segmentation using deep learning. *Acta Astronautica* 249 (2026), 381–398. https://doi.org/10.1016/j.actaastro.2026.07.010

[10] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. 2015. U-Net: Convolutional Networks for Biomedical Image Segmentation. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*. 234–241.

[11] tushhaaarrrr. 2026. Solar Filament Segmentation 2026 (0.40). https://www.kaggle.com/code/tushhaaarrrr/solar-filament-segmentation-2026-0-40. Kaggle notebook.

[12] Zongwei Zhou, Md Mahfuzur Rahman Siddiquee, Nima Tajbakhsh, and Jianming Liang. 2018. UNet++: A Nested U-Net Architecture for Medical Image Segmentation. In *Deep Learning in Medical Image Analysis and Multimodal Learning for Clinical Decision Support (DLMIA)*. 3–11.